

Terraform Provisioners

Terraform provisioners are used to execute scripts or commands on a local machine or on a remote resource after it is created or before it's destroyed. While provisioners can be powerful, they are considered a *last resort* in Terraform best practices, which recommend using cloud-init, configuration management tools (like Ansible), or image baking instead.

Types of Provisioners in Terraform

Provisioners are executed in two phases:

- **Creation-time Provisioners:** Run after a resource is created.
- **Destroy-time Provisioners:** Run before a resource is destroyed.

There are **two main categories** of provisioners:

1. local-exec Provisioner

- Executes a command **on the machine running Terraform**.
- Useful for operations like sending a Slack notification, calling an API, or running a local script.

```
resource "aws_instance" "example" {  
  
    ami          = "ami-0c55b159cbfafa1f0"  
  
    instance_type = "t2.micro"  
  
    provisioner "local-exec" {  
  
        command = "echo ${self.private_ip} >> private_ips.txt"  
  
    }  
  
}
```

2. remote-exec Provisioner

- **Connects to the provisioned resource (like EC2) via SSH or WinRM and runs commands.**
- **Requires credentials or connection block.**

```
resource "aws_instance" "example" {
```

```
ami          = "ami-0c55b159cbfafa1f0"

instance_type = "t2.micro"

key_name     = "my-key"


connection {

  type      = "ssh"

  user      = "ec2-user"

  private_key = file("~/ssh/my-key.pem")

  host      = self.public_ip

}
```

```
provisioner "remote-exec" {

  inline = [

    "sudo yum update -y",

    "sudo yum install -y nginx",

    "sudo systemctl start nginx"

  ]

}
```

Common Settings for Provisioners

- **inline:** Run multiple commands as an array.
- **script:** Reference to a local script file.
- **scripts:** Array of script file paths (deprecated in newer versions).
- **environment:** Define environment variables to use with remote-exec.

Destroy-Time Provisioners

You can specify a `when = "destroy"` block to run commands **before** resource destruction:

```
provisioner "remote-exec" {
```

```
when    = "destroy"

inline = ["echo 'Resource is being destroyed'"]

}
```

Connection Block

Used by remote-exec or file provisioners to connect to the resource:

```
connection {

  type    = "ssh"

  user    = "ubuntu"

  private_key = file("~/ssh/id_rsa")

  host     = self.public_ip

}
```

File Provisioner

Used to **upload files** from the local system to the remote resource.

```
provisioner "file" {

  source     = "app.conf"

  destination = "/etc/myapp/app.conf"

}
```

Provisioner Failure Behavior

By default, if a provisioner fails, Terraform marks the resource as **tainted**, and it will be destroyed and recreated in the next apply.

You can control this with:

- `on_failure = "continue"` → Terraform will not stop even if provisioner fails.

```
provisioner "remote-exec" {

  on_failure = "continue"

  inline    = ["some command"]

}
```

Best Practices and Warnings

- **Avoid provisioners when possible:** Use baked AMIs or cloud-init for initialization.
- **Provisioners can introduce non-determinism:** Commands may fail if the remote host isn't ready yet.
- **State Management Issues:** Terraform doesn't track what provisioners do, so re-running them requires destroying and recreating resources.
- **Debug with TF_LOG=DEBUG:** Helps when provisioners fail unexpectedly.

⚠ **Note:** Use provisioners only as a last resort. They are not idempotent, can cause unpredictable behavior, and are not recommended for regular infrastructure automation.